Multi-Layer Perceptrons

Johanni Brea

Introduction to Machine Learning

EPFL BIO322 2021

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Feature Engineering is great, but couldn't it be automatized?

With smart features basically any problem can be solved by a linear method.

How should we find the smart features?

Idea: Let us take a more flexible function family and find "features" and "regression coefficients" at the same time with gradient descent.

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Table of Contents

EPFL

Solving XOR
●○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

XOR-Problem
Training Data
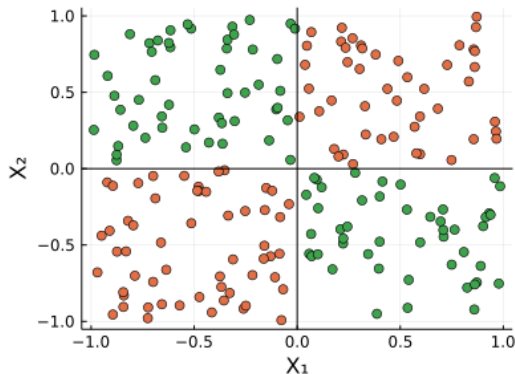
Logistic Regression fails:

There is no linear decision boundary.

EPFL

Solving XOR
○●○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Recap: Vector-Features

Project data to a higher dimensional space by computing the scalar products between feature vectors $w_1, \ldots, w_q$ and input vectors x_i and thresholding.

For example $h_{21} = \max(0, w_1^T x_2)$.

Logistic Regression on the features works.

Solving the XOR Problem without Feature Engineering

Logistic Regression: $P(Y = 1|x, \beta) = \sigma(\beta_0 + \beta_1 x_1 + \cdots \beta_p x_p)$

Logistic Regression on features:
$$P(Y = 1|x, \beta) = \sigma\big(\beta_0 + \beta_1 \underbrace{g(w_1^T x)}_{h_1} + \cdots \beta_q \underbrace{g(w_q^T x)}_{h_q}\big)$$

with hand-picked feature vectors $w_1, \ldots, w_q$ and activation function
$$g(x) = \text{relu}(x) = \max(0, x).$$

Idea

Why don't we learn the features with gradient descent?
$$P(Y = 1|x, \beta, \textcolor{red}{w_1, \ldots, w_q}) = \sigma\big(\beta_0 + \beta_1 g(w_1^T x) + \cdots \beta_q g(w_q^T x)\big)$$

$$\Rightarrow \quad \hat{\beta}, \hat{w} = \arg \min_{\beta, w_1, \ldots, w_q} \sum_{i=1}^{n} \log P(y_i|x_i, \beta, w_1, \ldots, w_q)$$

EPFL

Solving XOR
○○○●○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Solving the XOR Problem without Feature Engineering

It also works with learned features

- ▶ We just fitted our first neural network ☺.

- ▶ The loss function has local minima; gradient descent does not find for all initial guesses a good solution.

- ▶ With more than 4 feature vectors, gradient descent finds good solutions for most initial guesses.

Table of Contents

Artificial Neurons

Artificial neurons take a d-dimensional input $x = (x_1, \ldots, x_d)^T$ and output a scalar

$$a = g(w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_d x_d)$$

with **parameters (or weights)** $w_0, w_1, \ldots, w_d$ and **activation function** g.
w_0 is also called **bias** (instead of intercept).

EPFL

Solving XOR
○○○○○

Artificial Neurons
○●○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Popular Activation Functions

rectified linear unit $\quad \mathtt{relu}(x) = \max(0, x) = \begin{cases} x & x \geq 0, \\ 0 & x < 0 \end{cases}$

sigmoid $\quad \sigma(x) = \frac{1}{1+e^{-x}}$

tangent hyperbolic $\quad \tanh(x)$

softplus $\quad \mathtt{softplus}(x) = \log(\exp(x) + 1)$

heaviside (perceptron) $\quad H(x) = \begin{cases} 1 & x \geq 0, \\ 0 & x < 0 \end{cases}$

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○●○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Artificial relu-Neurons

$$a = \text{relu}(w_0 + w_1 x)$$

$$a = \text{relu}(w_0 + w_1 x_1 + w_2 x_2)$$

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○●

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
○○○

Table of Contents
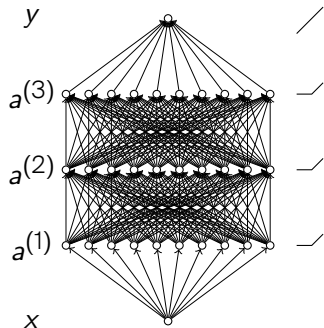
Multilayer Perceptrons

Multilayer Perceptrons (MLP) consist of multiple neurons organized in layers $1, 2, \ldots, L$.
Each **layer** has $d^{(l)}$ neurons and activation $g^{(l)}$.

output $a_k^{(l)}$ of k-th neuron in l-th layer

$$a_k^{(l)} = g^{(l)}\left(w_{k0}^{(l)} + w_{k1}^{(l)} a_1^{(l-1)} + \cdots + w_{kd^{(l-1)}}^{(l)} a_{d^{(l-1)}}^{(l-1)} \right)$$

input layer $a_k^{(0)} = x_k$.

matrix notation $a^{(l)} = g^{(l)}\left(w^{(l)} a^{(l-1)} + b^{(l)} \right)$
with $b_k^{(l)} = w_{k0}^{(l)}$.

one **input neuron** x
one linear **output neuron** y
3 **hidden layers** of 10 **relu-neurons**

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○●○○○

Regression
○○○○○○○

Classification
○○○

Multilayer Perceptrons
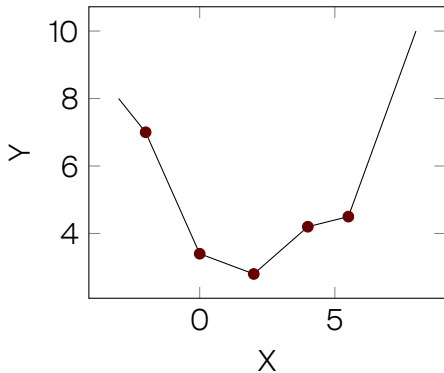
$g^{(1)} = $ relu, $g^{(1)} = $ tanh, $g^{(2)} = $ identity

2D input, $g^{(1)} = $ relu, $g^{(2)} = $ identity

Depth versus Width

1 hidden layer with 5 neurons

$10 + 6 = 16$ parameters

two hidden layers with 3 and 2 neurons

$6 + 8 + 3 = 17$ parameters

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○●○

Regression
○○○○○○○

Classification
○○○

Quiz

▶ The number of input units in a neural network is equal to the number of training samples.

▶ The activation of an artificial neuron with inputs $x_1 = 1$, $x_2 = 3$, $x_3 = 0$, weights $w_1 = 1$, $w_2 = -1$, $w_3 = 10$, bias $w_0 = 1$ and relu activation function is
A -1   **B** 0   **C** 1

▶ For a network with 3-dimensional input, 2 hidden layers of each 10 neurons and one output neuron the number of free parameters (weights and biases) is
A 24   **B** 131   **C** 161

▶ The XOR problem could also be solved with $g^{(1)}(x) = x$ and $g^{(2)}(x) = x$ (instead of $g^{(1)}(x) = \text{relu}(x)$ and $g^{(2)}(x) = x$).

Table of Contents

Regression with Multilayer Perceptrons

The output of the neural network is used to parametrize the conditional density.
The parameters are fitted with gradient descent
on the negative log-loglikelihood loss

$$\log \ell(\theta) = \sum_{i=1}^{n} \log p(y_i | x_i, \theta) \,.$$

For example: Assume the wind speed in Luzern is distributed normally around some mean
that correlates with the measurements done 5 hours earlier.

We take a neural network with as many input neurons as measurements, some hidden
neurons and one output neuron. Gradient descent finds the parameters such that the
output activity approaches the mean of the conditional density.

$$p(y|x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2} \left(y - g^{(2)}(b^{(2)} + w^{(2)}g^{(1)}(b^{(1)} + w^{(1)}x)) \right)^2 \right)$$

Regression with Multilayer Perceptrons

A neural network with more outputs can be used to predict more complex densities.

Example: Assume the wind speed in Luzern is log-normally distributed with mean and standard deviation correlating with the pressure in Luzern.

We take a neural network with one input neuron, some hidden neurons and two output neurons. Gradient descent finds the parameters such that the output neurons code for the mean and the variance[1] of the log-normal density.

$$p(y|x) = \frac{1}{\sqrt{2\pi f\left(a_2^{(2)}\right)}} \exp\left(-\frac{1}{2f\left(a_2^{(2)}\right)} \left(\log(y) - a_1^{(2)}\right)^2\right)$$

---

[1] The output of the second neuron is additionally transformed with function f to be positive.

Initialization and Data Preprocessing

The initialization of neural networks and the choice the hyper-parameters (number of hidden neurons, non-linearities, learning rate, etc.) is an art.

Common choices are:

biases = 0

weights $w_{ij}^{(l)}$ sampled uniformly from $[-x, x]$ with $x = \sqrt{\frac{6}{d^{(l-1)} + d^{(l)}}}$

learning rate between 10^{-4} and 10^{-1}.

These choices work typically well for input data between 0 and 1 or standardized input with each predictor having mean 0 and variance 1.

For regression it is advisable to also scale or standardize the output.

Regularization

For the weights and biases in each layer one can apply L1 or L2 regularization.
Early stopping in gradient descent can be used.
Using fewer hidden neurons also reduces the flexibility of the neural network.

Another popular and effective regularization method is **Dropout**:
During training, for each training example a randomly selected fraction of p neurons
is dropped out (inactivated).

This prevents neurons from becoming over-specialized.
All neurons are active when testing, but their weights are scaled by $1 - p$.

Flexibility of a Neural Network and Its Number of Parameters

More layers or more neurons $\Rightarrow$ more parameters.

More parameters $\Rightarrow$ more flexibility?

Not necessarily! Regularization has a strong effect on the flexibility.
Even without explicit regularization (L1, L2, Dropout) and without explicitly monitored early stopping one does stop gradient descent usually after some number of iterations and before perfect convergence; therefore one regularizes by implicit early stopping.

Wisely regularized large neural networks often work better than small ones, because they tend not to get stuck at sub-optimal losses.

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○●○

Classification
○○○

Flexibility by Composition of Simple Elements.

▶ Individual neurons should not be simpler: the composition of linear functions is a linear function.

▶ Individual neurons do not need to be more complex: complexity is achieved by using multiple neurons.

▶ With sufficiently many neurons one can approximate any function.

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○●

Classification
○○○

Table of Contents

EPFL

Solving XOR
○○○○○

Artificial Neurons
○○○○

Multilayer Perceptrons
○○○○○

Regression
○○○○○○○

Classification
●○○

Classification with Multilayer Perceptrons

With K output neurons we can use neural networks to parametrize categorical distributions suitable for classification problems.

Example: We take $28 \times 28 = 784$ input neurons, some hidden neurons and 10 output neurons to classify MNIST images. The softmax of the 10 output neurons is the predicted probability of the different class labels.

$$P(C_i|x) = s\left(g^{(2)}(b^{(2)} + w^{(2)}g^{(1)}(b^{(1)} + w^{(1)}x))\right)$$

where s is the softmax function (see slides "Supervised Learning").

Quiz

- ▶ With L1 or L2 regularization applied to the weights of a neural network, the final weights at the end of gradient descent tend to be smaller than without regularization.
- ▶ With early stopping gradient descent usually stops in a local minimum of the training loss.
- ▶ A neural network with no hidden layers and sigmoid activation function is equivalent to logistic regression.
- ▶ Gradient descent in neural networks always finds the global minimum of the loss function of the training set.
- ▶ Which activation function should be chosen in the output layer to predict the mean in a regression setting?

  A relu   **B** tanh   **C** identity